

CC LAB PRGRMS

2. Write a C program for storing data in a simulated cloud storage environment using a local file.

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    char data[1024];

    // Open a file in write mode. Simulating a cloud storage by using a local file.
    FILE *fptr = fopen("cloud_storage.txt", "w");
    if (fptr == NULL) {
        printf("Error opening the file!\n");
        exit(1);
    }

    // Get user input
    printf("Enter text to store in the cloud: ");
    fgets(data, sizeof(data), stdin);

    // Write data to the file
    fprintf(fptr, "%s", data);

    // Close the file
    fclose(fptr);

    printf("Data successfully saved to 'cloud_storage.txt'\n");

    return 0;
}
```

3. Write a C-Program for CPU usage Monitoring and Logging on Cloud-Based Ubuntu Server.

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <string.h>

// Function to read the current CPU usage from the /proc/stat file
float get_cpu_usage() {
    long double a[4], b[4], loadavg;
    FILE *fp;

    // Read first set of CPU statistics
```

```

fp = fopen("/proc/stat","r");
if (fp == NULL) {
    perror("Failed to open /proc/stat");
    exit(EXIT_FAILURE);
}
fscanf(fp, "%*s %Lf %Lf %Lf %Lf", &a[0], &a[1], &a[2], &a[3]);
fclose(fp);

// Sleep for a second to get the second set of CPU statistics
sleep(1);

// Read second set of CPU statistics
fp = fopen("/proc/stat","r");
if (fp == NULL) {
    perror("Failed to open /proc/stat");
    exit(EXIT_FAILURE);
}
fscanf(fp, "%*s %Lf %Lf %Lf %Lf", &b[0], &b[1], &b[2], &b[3]);
fclose(fp);

// Calculate the CPU usage
loadavg = ((b[0]+b[1]+b[2]) - (a[0]+a[1]+a[2])) /
           ((b[0]+b[1]+b[2]+b[3]) - (a[0]+a[1]+a[2]+a[3]));

return loadavg;
}

int main() {
    FILE *log_file;
    char *filename = "cpu_usage.log";

    // Open log file for writing
    log_file = fopen(filename, "w");
    if (log_file == NULL) {
        perror("Failed to open log file");
        return EXIT_FAILURE;
    }

    // Loop to record CPU usage
    for (int i = 0; i < 10; ++i) {
        float usage = get_cpu_usage();
        fprintf(log_file, "CPU Usage: %.2f%%\n", usage * 100);
        printf("Logged CPU Usage: %.2f%%\n", usage * 100);
    }
}

```

```

    // Sleep for 5 seconds
    sleep(5);
}

// Close log file
fclose(log_file);
printf("CPU usage logging completed.\n");

return 0;
}

```

4. Use Google App Engine Launcher to launch the web applications.

Step 1: Install Google App Engine

1. Install Google App Engine 1.9.62:

Go To: <https://www.npackd.org/p/com.google.AppEnginePythonSDK/1.9.62> and install the 1.9.62 version.

- Go to [npackd.org](https://www.npackd.org) for Google App Engine Python SDK 1.9.62 and download the 1.9.62 version.
- Ensure Python 2.7 is installed on your system. If not, download and install it from python.org.

Step 2: Set Up Google App Engine Launcher

1. Create a New Folder on Desktop:

- Right-click on the desktop, select **New -> Folder**, and name it appropriately (e.g., **MyGAEApp**).

2. Configure Preferences:

- Open Google App Engine Launcher.
- Go to **Edit -> Preferences**.
- Set the path for Google App Engine and Python.
 - **Google App Engine Path:** This should be the installation directory of the Google App Engine SDK.
 - **Python Path:** This should be the path where Python 2.7 is installed (typically **C:\Python27** on Windows).

Step 3: Create a New Application

1. Create a New Application:

- In Google App Engine Launcher, go to **File -> Create New Application**.

- Select the path of the folder created earlier (e.g., `MyGAEApp`).
- 2. **Configure the Application Folder:**
 - Open the created folder. You should see an automatically created sub-folder (e.g., `MyGAEApp`).
 - Ensure the project directory contains an `app.yaml` file. This file provides instructions for Google App Engine to provision resources for your app.

Example `app.yaml` File

Create an `app.yaml` file in the project directory with the following content:

Example Python Application (`main.py`)

Create a `main.py` file in the project directory with the following content:

```
import webapp2

class MainPage(webapp2.RequestHandler):
    def get(self):
        self.response.headers['Content-Type'] = 'text/plain'
        self.response.write('Hello, Google App Engine!')

app = webapp2.WSGIApplication([
    ('/', MainPage),
], debug=True)
```

Step 4: Run the Application

1. **Launch the Application:**
 - In Google App Engine Launcher, select the created application.
 - Click the `Run` button.
 - The application will start running locally, and you should see output indicating that the server is running.
2. **Access the Application:**
 - Open a web browser and go to `http://localhost:8080`.
 - You should see the message `Hello, Google App Engine!`.

8. Fifa

Import Necessary Libraries:

- Import the required PySpark modules for creating a Spark session and defining the schema.

Initialize SparkSession:

- Create a Spark session with the application name "FIFADatasetExample".

Define the Schema for the FIFA Dataset:

- Define a schema using `StructType` and `StructField` to specify the data types for each column in the dataset.

Load FIFA Dataset from CSV:

- Read the CSV file into a DataFrame using the defined schema.

Show DataFrame:

- Display the contents of the DataFrame using `show()`.

Select and Show Specific Columns:

- Select the `previous_points` column from the DataFrame and display it.

Stop SparkSession:

- Stop the Spark session to free up resources.

Import necessary libraries

```
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, StringType, IntegerType, FloatType
```

Initialize SparkSession

```
spark = SparkSession.builder \
    .appName("FIFADatasetExample") \
    .getOrCreate()
```

Define the schema for the FIFA dataset

```
schema = StructType([
    StructField("team", StringType(), True),
    StructField("team_code", StringType(), True),
    StructField("association", StringType(), True),
    StructField("rank", IntegerType(), True),
    StructField("previous_rank", IntegerType(), True),
```

```

        StructField("points", FloatType(), True),
        StructField("previous_points", FloatType(), True)
    ])

# Load FIFA dataset from CSV
file_path = "/content/fifa_ranking_2022-10-06.csv"
fifa_df = spark.read.csv(file_path, header=True, schema=schema)

# Show DataFrame
fifa_df.show()

# Select and show specific columns
fifa_df.select("previous_points").show()

# Stop SparkSession
spark.stop()

```

5. Setting up a single-node Hadoop cluster involves several steps. Below is a structured and detailed guide with explanations and corrections to ensure clarity and completeness:

Step-by-Step Guide to Setting Up a Single-Node Hadoop Cluster

1. Update and Install Necessary Packages

```

```sh
sudo apt update
sudo apt install ssh pdsh
```

```

2. Configure SSH for Passwordless Login

```

```sh
ssh-keygen -t rsa -P ""
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
ssh localhost
```

```

3. Install Java (Hadoop Requires Java)

```

```sh
sudo apt install openjdk-8-jdk
java -version
```

```

4. Download and Extract Hadoop

```

```sh

```

```
sudo wget -P ~
https://archive.apache.org/dist/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz
tar xzf ~/hadoop-3.2.1.tar.gz
mv ~/hadoop-3.2.1 ~/hadoop
...
```

#### #### 5. Configure Hadoop Environment Variables

Edit `hadoop-env.sh` to set the `JAVA\_HOME` path:

```
```sh
nano ~/hadoop/etc/hadoop/hadoop-env.sh
...
```

Add the following line:

```
```sh
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/
...
```

#### #### 6. Move Hadoop Directory to `/usr/local` and Set Permissions

```
```sh
sudo mv ~/hadoop /usr/local/hadoop
sudo chown -R $(whoami):$(whoami) /usr/local/hadoop
sudo chmod -R 755 /usr/local/hadoop
...
```

7. Update System Environment Variables

Edit `/etc/environment`:

```
```sh
sudo nano /etc/environment
...
```

Add the following lines:

```
```sh
HADOOP_HOME="/usr/local/hadoop"
PATH="$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin"
JAVA_HOME="/usr/lib/jvm/java-8-openjdk-amd64"
...
```

Reload the environment variables:

```
```sh
source /etc/environment
...
```

#### #### 8. Create a Hadoop User (Optional)

```
```sh
sudo adduser hadoopuser
sudo usermod -aG sudo hadoopuser
...
```

9. Configure Hadoop Core and HDFS

Edit `core-site.xml`:

```
```sh
sudo nano /usr/local/hadoop/etc/hadoop/core-site.xml
```
```

Add the following configuration:

```
```xml
<configuration>
 <property>
 <name>fs.defaultFS</name>
 <value>hdfs://localhost:9000</value>
 </property>
 <property>
 <name>hadoop.tmp.dir</name>
 <value>/usr/local/hadoop/tmp</value>
 </property>
</configuration>
```
```

Edit `hdfs-site.xml`:

```
```sh
sudo nano /usr/local/hadoop/etc/hadoop/hdfs-site.xml
```
```

Add the following configuration:

```
```xml
<configuration>
 <property>
 <name>dfs.replication</name>
 <value>1</value>
 </property>
</configuration>
```
```

10. Configure MapReduce and YARN

Edit `mapred-site.xml`:

```
```sh
sudo nano /usr/local/hadoop/etc/hadoop/mapred-site.xml
```
```

Add the following configuration:

```
```xml
<configuration>
 <property>
 <name>mapreduce.framework.name</name>
 </property>
</configuration>
```
```



```
        <value>yarn</value>
    </property>
</configuration>
...
```

Edit `yarn-site.xml`:

```
```sh
sudo nano /usr/local/hadoop/etc/hadoop/yarn-site.xml
...
```

Add the following configuration:

```
```xml
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
  </property>
</configuration>
...
```

11. Format the HDFS Filesystem

```
```sh
hdfs namenode -format
...
```

#### #### 12. Start Hadoop Services

```
```sh
start-dfs.sh
start-yarn.sh
...
```

13. Verify Hadoop Installation

You can check the Hadoop services running by accessing the web UIs:

- NameNode: http://localhost:9870
- ResourceManager: http://localhost:8088

14. Clean Up (Optional)

Remove `pdsh` if no longer needed:

```
```sh
sudo apt remove pdsh
```

## 6.Hadoop Programming: Word Count Map Reduce Program Using Eclipse.

### 1. Driver Code (`WCDriver.java`)

The driver code configures the Hadoop job and sets up the mapper and reducer classes.

```
``java
import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WCDriver extends Configured implements Tool {
 public int run(String[] args) throws IOException {
 if (args.length < 2) {
 System.out.println("Please provide valid inputs");
 return -1;
 }
 JobConf conf = new JobConf(WCDriver.class);
 FileInputFormat.setInputPaths(conf, new Path(args[0]));
 FileOutputFormat.setOutputPath(conf, new Path(args[1]));
 conf.setMapperClass(WCMapper.class);
 conf.setReducerClass(WCReducer.class);
 conf.setMapOutputKeyClass(Text.class);
 conf.setMapOutputValueClass(IntWritable.class);
 conf.setOutputKeyClass(Text.class);
 conf.setOutputValueClass(IntWritable.class);
 JobClient.runJob(conf);
 return 0;
 }

 public static void main(String[] args) throws Exception {
 int exitCode = ToolRunner.run(new WCDriver(), args);
 System.out.println(exitCode);
 }
}
```

#### #### 2. Mapper Code (`WCMapper.java`)

The mapper code processes each line of input and emits key-value pairs for each word.

```
```java
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable,
Text, Text, IntWritable> {
    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
output, Reporter rep) throws IOException {
        String line = value.toString();
        // Splitting the line on spaces
        for (String word : line.split(" ")) {
            if (word.length() > 0) {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}
```
```

#### #### 3. Reducer Code (`WCReducer.java`)

The reducer code aggregates the counts for each word and emits the final count.

```
```java
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text,
IntWritable, Text, IntWritable> {
```

```

    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text,
IntWritable> output, Reporter rep) throws IOException {
        int count = 0;
        // Counting the frequency of each word
        while (values.hasNext()) {
            count += values.next().get();
        }
        output.collect(key, new IntWritable(count));
    }
}
...

```

Compiling and Running the MapReduce Job

1. ****Compile the Java Code**:**

- Save the three Java files (`WCDriver.java`, `WCMapper.java`, `WCReducer.java`) in the same directory.
- Open a terminal and navigate to the directory containing these files.
- Compile the code using the following commands:

```

```sh
javac -classpath $(hadoop classpath) -d . WCDriver.java WCMapper.java
WCReducer.java
...

```

#### 2. **\*\*Create a JAR File\*\*:**

- Package the compiled classes into a JAR file:

```

```sh
jar cvf hadoop-mapreduce-example.jar *.class
...

```

3. ****Run the MapReduce Job**:**

- Ensure your Hadoop cluster is running.
- Copy your input data to HDFS:

```

```sh
hdfs dfs -mkdir /sample
hdfs dfs -mkdir /sample/input
hdfs dfs -put /path/to/your/input/file /sample/input
...

```

- Run the Hadoop job using the following command:

```
```sh
hadoop jar hadoop-mapreduce-example.jar WCDriver /sample/input /sample/output
```
```

#### 4. **\*\*Check the Output\*\*:**

- After the job completes, check the output in HDFS:

```
```sh
hdfs dfs -cat /sample/output/part-00000
```
```

### 7. File Management tasks in Hadoop using HDFS commands

- a. Adding files to HDFS
- b. Retrieving files from HDFS
- c. Deleting files from HDFS

#### **## File Management Tasks in Hadoop Using HDFS Commands**

##### **### 1. Create a Directory in HDFS**

###### **#### Usage:**

```
```sh
hadoop fs -mkdir <paths>
```
```

###### **#### Example:**

```
```sh
hadoop fs -mkdir /user/saurzcode/dir1 /user/saurzcode/dir2
```
```

##### **### 2. List the Contents of a Directory**

###### **#### Usage:**

```
```sh
hadoop fs -ls <args>
```
```

###### **#### Example:**

```
```sh
hadoop fs -ls /user/saurzcode
```
```

##### **### 3. Upload and Download a File in HDFS**

###### **#### a. Upload a File to HDFS**

**\*\*Usage:\*\***

```
```sh
hadoop fs -put <localsrc> ... <HDFS_dest_Path>
```
```

**\*\*Example:\*\***

```
```sh
hadoop fs -put /home/saurzcode/Samplefile.txt /user/saurzcode/dir3/
```
```

#### #### b. Download a File from HDFS

**\*\*Usage:\*\***

```
```sh
hadoop fs -get <hdfs_src> <localdst>
```
```

**\*\*Example:\*\***

```
```sh
hadoop fs -get /user/saurzcode/dir3/Samplefile.txt /home/
```
```

#### ### 4. See the Contents of a File

**\*\*Usage:\*\***

```
```sh
hadoop fs -cat <path[filename]>
```
```

**\*\*Example:\*\***

```
```sh
hadoop fs -cat /user/saurzcode/dir1/abc.txt
```
```

#### ### 5. Remove a File or Directory in HDFS

**\*\*Usage:\*\***

```
```sh
hadoop fs -rm <arg>
```
```

**\*\*Example:\*\***

```
```sh
```

```
hadoop fs -rm /user/saurzcode/dir1/abc.txt
...

```

6. Display Last Few Lines of a File

****Usage:****

```
```sh
hadoop fs -tail <path[filename]>
...

```

**\*\*Example:\*\***

```
```sh
hadoop fs -tail /user/saurzcode/dir/abc.txt
...

```

Summary of Commands

Task	Command
Create a directory	<code>`hadoop fs -mkdir <paths>`</code>
List directory contents	<code>`hadoop fs -ls <args>`</code>
Upload a file to HDFS	<code>`hadoop fs -put <localsrc> ... <HDFS_dest_Path>`</code>
Download a file from HDFS	<code>`hadoop fs -get <hdfs_src> <localdst>`</code>
See contents of a file	<code>`hadoop fs -cat <path[filename]>`</code>
Remove a file or directory	<code>`hadoop fs -rm <arg>`</code>
Display last few lines of a file	<code>`hadoop fs -tail <path[filename]>`</code>